


GitHub Repo:

<https://github.com/Muntasir101/Assessment-5-LoRAandQLoRA>

Output:




```
[ ] # Train the model
trainer.train()

Map: 0%|          | 0/4 [00:00<?, ? examples/s]
Map: 0%|          | 0/1 [00:00<?, ? examples/s]
Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased and are newly i
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
/tmp/ipython-input-362689685.py:24: FutureWarning: `tokenizer` is deprecated and will be removed in version 5.0.0 for `Trainer.__init__`. Use
trainer = Trainer(
/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:666: UserWarning: 'pin_memory' argument is set as true but no accelera
warnings.warn(warn_msg)

[6/6 00:42, Epoch 3/3]

Epoch   Training Loss   Validation Loss
-----
1         No log           0.726608
2         No log           0.720246
3         No log           0.720030

/usr/local/lib/python3.12/dist-packages/torch/utils/data/dataloader.py:666: UserWarning: 'pin_memory' argument is set as true but no accelera
warnings.warn(warn_msg)
Training Loss: 0.720030, Validation Loss: 0.720030, Predicted Sentiment: Positive
```



```
[ ] # Apply the example to the model
with torch.no_grad():
    outputs = model(**inputs)

# Extract the new predictions
predicted_labels = torch.argmax(outputs.logits, dim=1).tolist()

sentiment_map = {0: "Negative", 1: "Positive"}

# Since the input text is a single example, we take the first prediction
predicted_sentiment = sentiment_map[predicted_labels[0]]

print(f"Predicted Sentiment: {predicted_sentiment}")

Predicted Sentiment: Positive
```